



CS2208: Information Storage and Management I

Dr. Alejandro Arbelaez



Dr. Alejandro Arbelaez

- Lecturer – Computer Science Department
- Experience:
 - +4 Years teaching experience at third level (CS - Data Science & Analytics)
 - +4 Postdoctoral Researcher
 - PhD. In Computer Science
 - Engineering degree in Computer Science



 a.arbelaez@cs.ucc.ie

Dr. Alejandro Arbelaez



Dr. Alejandro Arbelaez

- 2019 – Present: Lecturer at UCC
- 2017 – 2019: Lecturer at CIT
 - Teaching mainly in the MSc in AI, MSc in Data Analytics, and Software Development programme
- 2013 – 2017: Senior Postdoc at University College Cork
 - Working in Data Science and Analytics projects
- 2011- 2013: Postdoc at University of Tokyo / Franco-Japanese Research Lab
 - Working in AI and Massively parallel computing
- 2007 - 2011 : PhD Candidate at Université Paris XI
 - Working in the interception between descriptive predictive and prescriptive analytics

Learning Outcomes

Module Description

- Design relational databases for a range of data types;
- Demonstrate a detailed knowledge of the SQL language and SQL-based database management systems;
- Demonstrate a working knowledge of the principles and practices of relational database design and administration;
- Apply database management principles to real-world application domains, such as biology, business, and science.

Module Description

CS2208

CS2208	Information Storage and Management I	Computer Science	5	- Module Description
Credit Weighting: 5				
Semester(s): Semester 1.				
No. of Students: Max 30.				
Pre-requisite(s): CS1106				
Co-requisite(s): None				
Teaching Method(s): 24 x 1hr(s) Lectures; 10hr(s) Practicals.				
Module Co-ordinator: Dr Alejandro Arbelaez, Department of Computer Science.				
Lecturer(s): Dr Alejandro Arbelaez, Department of Computer Science.				
Module Objective: Students will learn: analysis requirements for various types of application for managing persistent data and how to design, implement and administer databases to meet these requirements; the remainder of the SQL concepts and constructs not covered in the prerequisite module.				
Module Content: Database Management Systems; DBMS storage structures and indexing. Relational algebra and relational calculus; SQL; query optimisation; views. Database Design: conceptual, logical and physical database design; Keys; data integrity; functional dependencies and normal forms; Object-relational databases; Database triggers.				
Learning Outcomes: On successful completion of this module, students should be able to: <ul style="list-style-type: none">▪ design relational databases for a range of data types;▪ demonstrate a detailed knowledge of the SQL language and SQL-based database management systems;▪ demonstrate a working knowledge of the principles and practices of relational database design and administration;▪ Apply database management principles to real-world application domains, such as biology, business, and science.				
Assessment: Total Marks 100: Formal Written Examination 80 marks; Continuous Assessment 20 marks (Assignments and/or in-class tests).				
Compulsory Elements: Formal Written Examination; Continuous Assessment.				

Note: Every effort has been made to ensure that the programme and module content as described in the University's Calendar and Book of Modules for the 2020-21 academic year are accurate. However, due to the Coronavirus (COVID-19) pandemic, no guarantee is given that programme/module content, delivery and assessment may not be altered, cancelled, replaced, augmented or otherwise amended. Any changes will ensure the same competencies and Learning Outcomes are met. Programme and/or Module Coordinators will communicate any such changes to students.

Module workload

- 2 hours lecture every week (24 hours)
- 1 hours weekly Lab

Assessment Breakdown

- This module includes a combination of Continuous Assessment and a Written Exam
 - Continuous Assessment: 20%
 - Interactive Canvas exam (via Canvas quizzes) 10% - Week 6
 - Project Assignment 10% - End of semester (approx. Week 11)
 - Written Exam: 80%

Canvas

Plagiarism

1. Plagiarism is presenting someone else's work as your own. It is a violation of UCC Policy and there are strict and severe penalties.
2. You must read and comply with the UCC Policy on Plagiarism www.ucc.ie/en/exams/procedures-regulations/
3. The Policy applies to *all* work submitted, including software.
4. You can expect that your work will be checked for evidence of plagiarism or collusion.
5. In some circumstances it may be acceptable to reuse a small amount of work by others, but *only* if you provide explicit acknowledgement and justification.
6. If in doubt ask your module lecturer *prior* to submission. Better safe than sorry!

This Module

Basic SQL

- Create tables, basic join, primary keys, etc


DB Modelling

- ER Diagrams, Normal Forms, Relational Algebra/calculus

Advanced SQL


- Indexing, views, triggers, etc.

Resources → Books



Carolyn E. Begg and Thomas M. Connolly, **Database Systems A Practical Approach to Design, Implementation, and Management** [ISBN-10: 0321523067]

A. Silberschatz, H. F. Korth, S. Sudarshanm **Database System Concepts**, McGraw-Hill



Databases

- What is a database?
 - A collection of files storing related data
- Example of databases
 - Accounts databases
 - Payroll database
 - UCC's students database
 - Amazon's products databases
 - Airline reservation database

Database Management System

- What is a DBMS?
 - A big program written by someone else that allows us to manage efficiently a large database and allows it to persist over long periods of time
- Examples of DBMSs
 - Oracle, IBM DB2, Microsoft SQL Server, etc
 - Open source: MySQL, PostgreSQL, CouchDB
 - Open source library: SQLite
- This semester we will focus on relational DBMS

Example – Online Bookseller

- What data do we need?
 - Data about books, customers, pending orders, order histories, trends performances, etc.
 - Data about sessions (clicks, pages, search history)
 - Note: **data must be persistent!**
 - Also note that data is large.. Won't fit all in memory
- What capabilities on the data do we need?
 - Insert/remove books, find books by author/title/etc, analyze past order history, recommend books, etc.
 - Data must be accessed **efficiently** by many users
 - Data must be **safe from failures** and malicious users

Using Databases

- Jane and John both have a shared ID number for a gift (credit) of \$200 they got as a wedding gift
 - Jane @ her office orders “The selfish Gene”, \$80
 - John @ his office orders “Guns the Steel”, \$100
- Questions
 - What is the ending credit?
 - What if the second book costs \$130?
 - What if the system crashes?
- **A DBMS needs to handle various users issues!**

What functions should a DBMS provide?

1. Describe real-world entities in terms of stored data
2. Persistently store large datasets
3. Efficiently query & update
 - Must handle complex questions about data
 - Must handle sophisticated updates
 - Performance matters
4. Change structure (e.g., add attributes)
5. Concurrency control: enable simultaneous updates
6. Crash recovery
7. Security and integrity

DBMS Benefits

- Expensive to implement all these features inside the application
- DBMS provides these features (and more)
- DBMS simplifies application development

Key Data Management Concepts

- **Data models:** how to describe real-world data
 - Relational, NoSQL, etc..
- **Declarative query languages**
 - Say what you want not how to get it
- **Data independence**
 - Physical independence: can change how data is stored on disk without maintenance to applications
 - Logical independence: can change schema w/o affecting app
- **Query Optimizer**
 - Query plans and how they are executed
- **Physical design**
- **Transactions**
 - Isolation and atomicity

Relational Model

- Data is a collection of relations/tables:

Rows/
Tuples/
Records



cname	country	no_employees	for_profit
IBM	USA	20000	True
Sony	Japan	5000	True
Nintendo	Japan	3000	True
AirCanada	Canada	5000	True

- Mathematically, relation is a set of tuples
 - Each tuple appears 0 or 1 times in the table
 - Order the rows in unspecified

The Relational Data Model

- Degree or arity of a relation
 - Number of attributes
- Each attribute has a type
 - String: CHAR(20), VARCHAR(50), TEXT
 - Numbers: INT, FLOAT
 - Money, DateTime
 - Few more that are database specific
- **Statically and strictly enforced**

Keys

- An attribute that uniquely identifies a record
- A key can consists of multiple attributes
- Foreign key:
 - A attribute(s) that is a key for other relations

Relational Model: Example

Company(cname, country, no_employees, for_profit)

Country(name, population)

cname	country	no_employees	for_profit
IBM	USA	20000	True
Sony	Japan	5000	True
Nintendo	Japan	3000	True
AirCanada	Canada	5000	True

name	population
USA	320M
Japan	127M

Query Language


- SQL
 - **Structured Query Language**
 - Developed by IBM in the 70s
 - Most widely used language to query relational data
- We will see other languages for the relational model later on
 - Relational Algebra, Relational Calculus

/ˈsiːkwəl/

“SQL”

/ˈɛs kjuː ˈɛl/

1970s



/ˈsiːkwəl/

SEQUEL


Structured English **QUE**ry Language

1970s

SEQUEL

Structured English **Q**U**E**ry Language

/ˈsi:kwəl/



SQL

Structured **Q**uery Language

/ˈɛs kju: ˈɛl/